

Fearless Concurrency — Threads

Performant Software Systems with Rust — Lecture 12

Baochun Li, Professor

Department of Electrical and Computer Engineering
University of Toronto

Rust's Fearless Concurrency

- Rust's **ownership** and **type checking** features helped manage both memory safety and concurrency problems
- Many runtime concurrency errors becomes compile-time errors!
- Other programming languages, such as Erlang and Pony, may impose limits or performance tradeoffs
 - They implemented the **Actor Model**, which has elegant functionality for message-passing concurrency, but only obscure ways to share state between threads

Concurrent vs. Parallel Programming

- **Concurrent programming**: different parts of a program execute independently, but may not be at the same time
- **Parallel programming**: different parts of a program execute at the same time
- When we mention **concurrency**, we imply **concurrent or parallel** programming

Using Threads to Run Code Simultaneously

- **Multi-threading** allows the execution of multiple **threads** in parallel, at the same time
- But it also leads to subtle problems — a conventional topic in an OS course
 - Race conditions
 - Deadlocks
 - Thread synchronization
- Rust standard library uses a **1:1 (kernel threads)** model

Creating a New Thread with `spawn`

```
1 use std::thread;
2 use std::time::Duration;
3
4 fn main() {
5     // passing a closure to `spawn`
6     thread::spawn(|| {
7         for i in 1..10 {
8             println!("Printing number {i} from the spawned thread!");
9             thread::sleep(Duration::from_secs(1));
10        }
11    });
12
13    for i in 1..5 {
14        println!("Printing number {i} from the main thread!");
15        thread::sleep(Duration::from_secs(1));
16    }
17 }
```

Waiting for All Threads to Finish

```
1 fn main() {
2     // create a join handle to the spawned thread
3     let handle = thread::spawn(|| {
4         for i in 1..10 {
5             println!("hi number {i} from the spawned thread!");
6             thread::sleep(Duration::from_secs(1));
7         }
8     });
9
10    for i in 1..5 {
11        println!("hi number {i} from the main thread!");
12        thread::sleep(Duration::from_sec(1));
13    }
14
15    // wait for the spawned thread to finish
16    handle.join().unwrap();
17 }
```

Live Demo

Capturing the Environment of the Parent Thread

```
1 fn main() {
2     let v = vec![1, 2, 3];
3
4     let handle = thread::spawn(|| {
5         println!("Here's a vector: {v:?}");
6     });
7
8     handle.join().unwrap();
9 }
```

Live Demo

Capturing the Environment of the Parent Thread

```
1 fn main() {
2     let v = vec![1, 2, 3];
3
4     // `move` takes ownership of the environment
5     let handle = thread::spawn(move || {
6         println!("Here's a vector: {v:?}");
7     });
8
9     handle.join().unwrap();
10 }
```

Do not communicate by sharing memory; instead, share memory by communicating.

— The Go Language Documentation

The Actor Model

- An **actor** is the basic building block of concurrent computation
- In responding to messages that it receives, an **actor** makes local decisions, creates more actors, sends more messages, and modifies private states
- The **Actor Model** removes the need for lock-based synchronization
 - a mandatory design pattern in Erlang and Pony

Channels

- A general programming concept by which data is sent from one thread to another
- Two halves: one or more **transmitter(s)** and one or more **receiver(s)**
 - The transmitter half is the upstream location of a “river”, the receiver half is the downstream
 - **Closed** if either half is dropped

Multiple-Producer Single-Consumer Channels

```
1 use std::sync::mpsc;  
2  
3 fn main() {  
4     // returns a tuple that is destructured  
5     let (tx, rx) = mpsc::channel();  
6 }
```

send() / recv() vs. try_send() / try_recv()

- Both variants return `Result<T, E>`
- `send()` / `recv()` **blocks** the thread's execution and wait until a channel has available capacity or becomes non-empty
- `try_send()` / `try_recv()` is **non-blocking** and returns immediately

Implementing the Actor Model: Use channels for message passing

```
1 use std::sync::mpsc;
2 use std::thread;
3
4 fn main() {
5     // multiple producer, single consumer (MPSC)
6     let (tx, rx) = mpsc::channel();
7
8     thread::spawn(move || {
9         let val = String::from("hi");
10        tx.send(val).unwrap();
11    });
12
13    // use try_recv() to check if a message is available in
14    // a non-blocking way
15    let received = rx.recv().unwrap();
16    println!("Got: {}", received);
17 }
```

Transferring Ownership between Threads With Channels

Will this code compile successfully?

```
1 fn main() {
2     let (tx, rx) = mpsc::channel();
3
4     thread::spawn(move || {
5         let val = String::from("hi");
6         tx.send(val).unwrap();
7         println!("val is {}", val);
8     });
9
10    let received = rx.recv().unwrap();
11    println!("Got: {}", received);
12 }
```

Cloning Multiple Producers with an MPSC Channel

```
1 let (tx, rx) = mpsc::channel();
2
3 let tx1 = tx.clone();
4
5 thread::spawn(move || {
6     let vals = vec![
7         String::from("hi"),
8         String::from("from"),
9         String::from("the"),
10        String::from("thread"),
11    ];
12
13    for val in vals {
14        tx1.send(val).unwrap();
15        thread::sleep(Duration::from_secs(1));
16    }
17 });
18
```

**Again, do no communicate by sharing
memory, and keep the states in each
thread private and local!**

**But what if I really want to share
memory?**

Shared-State Concurrency with Mutex<T>

```
1 use std::sync::Mutex;
2
3 fn main() {
4     let m = Mutex::new(5);
5
6     {
7         // block the current thread until having the lock
8         // call to `lock` would fail if the holder thread panicked
9         // as `m` is Mutex<T>, cannot access its value directly
10        let mut num = m.lock().unwrap();
11        // returned value is a smart pointer, `MutexGuard`,
12        // which implements `Deref` and `Drop` traits
13        *num = 6;
14    }
15
16    println!("m = {:?}", m);
17 }
```

Sharing a Mutex<T> Between Threads

```
1 use std::sync::Mutex;
2 use std::thread;
3
4 fn main() {
5     let counter = Mutex::new(0);
6     let mut handles = vec![];
7
8     for _ in 0..10 {
9         let handle = thread::spawn(move || {
10             let mut num = counter.lock().unwrap();
11
12             *num += 1;
13         });
14         handles.push(handle);
15     }
16
17     for handle in handles {
18         handle.join().unwrap();
```

Will this compile successfully?

Multiple Ownership with Multiple Threads

```
1 use std::rc::Rc; // use reference counting to share ownership
2
3 fn main() {
4     let counter = Rc::new(Mutex::new(0));
5     let mut handles = vec![];
6
7     for _ in 0..10 {
8         let counter = Rc::clone(&counter);
9         let handle = thread::spawn(move || {
10             let mut num = counter.lock().unwrap();
11         // snip
12     }
13 }
```

Will this compile successfully?

Atomic Reference Counting with Arc<T>

```
1 use std::sync::{Arc, Mutex}; // use atomic reference counting instead
2 use std::thread;
3
4 fn main() {
5     let counter = Arc::new(Mutex::new(0));
6     let mut handles = vec![];
7
8     for _ in 0..10 {
9         let counter = Arc::clone(&counter);
10        let handle = thread::spawn(move || {
11            let mut num = counter.lock().unwrap();
```

Other Atomic Types in the Rust Standard Library

- Several atomic types that provide atomic access to primitive types
 - **safe** to share between threads
 - such as `AtomicUsize`, `AtomicBool`, and so on
- `load()` and `store()`
- Guaranteed to be **lock-free**
- Can be used as building blocks of other concurrent types

```
1 use std::sync::atomic::{AtomicUsize, Ordering};  
2  
3 static GLOBAL_THREAD_COUNT: AtomicUsize = AtomicUsize::new(0);  
4  
5 // relaxed ordering doesn't synchronize anything except the global  
6 // thread counter itself  
7 let old_thread_count = GLOBAL_THREAD_COUNT.fetch_add(1,  
8     Ordering::Relaxed);  
9  
10 // this number may not be true at the moment of printing because some  
11 // other thread may have changed static value already  
12 println!("live threads: {}", old_thread_count + 1);
```

Extensible Concurrency with the `Send` Marker Trait

- `Send`, a `std::marker` trait, indicates that ownership of values of the type implementing `Send` can be transferred between threads
- It is safe to **send** it to another thread
- Automatically implemented when the compiler thinks it's appropriate

Types That Are Send

- Almost every Rust type is `Send`, with a few exceptions
 - `Rc<T>` cannot be `Send` as both threads may update the reference count at the same time
- Any type composed entirely of `Send` types is automatically marked as `Send`

Extensible Concurrency with the Sync Marker Trait

- `Sync` indicates that it is safe for the type implementing `Sync` to be **referenced** from multiple threads
- A type `T` is `Sync` if and only if `&T` is `Send`
- Automatically implemented when the compiler thinks it's appropriate

Types That Are Sync

- Almost every Rust type is Sync, with a few exceptions
 - Rc<T> is not Sync
 - RefCell<T> is not Sync since its implementation of borrow checking at runtime is not thread-safe
 - Mutex<T> is Sync
- Any type composed entirely of Sync types is automatically marked as Sync

So long, multi-threading!

Required Additional Reading

The Rust Programming Language, Chapter 16 and 21