

Functional Rust

Performant Software Systems with Rust — Lecture 10

Baochun Li, Professor

Department of Electrical and Computer Engineering
University of Toronto

What Is Functional Programming?

- A programming paradigm where programs are constructed by **applying** and **composing** functions
- Functions are first-class citizens
 - bound to names
 - passed as arguments
 - returned from other functions
- Used in a wide variety of languages such has **Haskell** and **Clozure**

Closures

- Closures are **anonymous functions** that can be
 - saved in a variable
 - passed as arguments to other functions
- Closures **capture** values from the environment (or scope) in which they are defined

First Example

- Every so often, our t-shirt company gives away an exclusive, limited-edition shirt as a promotion
- People can optionally add their favorite color to their profile
- If the person chosen for a free shirt has their favorite color set, they get that color shirt
- Otherwise, they get whatever color the company currently has the most of

```
1 #[derive(Debug, Copy, Clone)]
2 enum ShirtColor {
3     Red,
4     Blue,
5 }
6
7 struct Inventory {
8     shirts: Vec<ShirtColor>,
9 }
```

```
1 impl Inventory {
2     fn most_stocked(&self) → ShirtColor {
3         let mut num_red = 0;
4         let mut num_blue = 0;
5
6         for color in &self.shirts {
7             match color {
8                 ShirtColor::Red => num_red += 1,
9                 ShirtColor::Blue => num_blue += 1,
10            }
11        }
12
13        if num_red > num_blue {
14            ShirtColor::Red
15        } else {
16            ShirtColor::Blue
17        }
18    }
}
```

```
1 impl Inventory {  
2     fn giveaway(&self, user_preference: Option<ShirtColor>)  
3         → ShirtColor {  
4             user_preference.unwrap_or_else(|| self.most_stocked())  
5         }  
6     }  
7 }
```

Closure Type Inference and Annotation

- Closures don't usually require you to annotate the types of the parameters or the return value
 - because they are not used in an exposed interface like `fn` functions do
 - rather, they are stored in variables and passed as arguments
- But we can add type annotations too, if needed

Closures with Type Annotation

```
1 let expensive_closure = |num: u32| → u32 {  
2     println!("calculating slowly...");  
3     thread::sleep(Duration::from_secs(2));  
4     num  
5 };
```

Closures with Type Annotation and Inference

```
1 fn add_one_v1 (x: u32) → u32 { x + 1 }
2 let add_one_v2 = |x: u32| → u32 { x + 1 };
3 let add_one_v3 = |x| { x + 1 };
4 let add_one_v4 = |x| x + 1;
5 // brackets optional: closure body has only one expression
```

What happens if we compile this code?

```
1 let example_closure = |x| x;  
2  
3 let s = example_closure(String::from("hello"));  
4 let n = example_closure(5);
```

Three Ways of Capturing Values From the Environment

- Borrowing immutably
- Borrowing mutably
- Taking ownership

Let's look at an example on each

Borrowing Immutably

```
1 fn main() {  
2     let list = vec![1, 2, 3];  
3     println!("Before defining closure: {list:?}");  
4  
5     let only_borrows = || println!("From closure: {list:?}");  
6  
7     println!("Before calling closure: {list:?}");  
8     only_borrows();  
9     println!("After calling closure: {list:?}");  
10 }
```

Borrowing Mutably

```
1 fn main() {  
2     let mut list = vec![1, 2, 3];  
3     println!("Before defining closure: {list:?}");  
4  
5     let mut borrows_mutably = || list.push(7);  
6  
7     borrows_mutably();  
8     println!("After calling closure: {list:?}");  
9 }
```

Taking Ownership

```
1 use std::thread;
2
3 fn main() {
4     let list = vec![1, 2, 3];
5     println!("Before defining closure: {list:?}");
6
7     thread::spawn(move || println!("From thread: {list:?}"))
8         .join()
9         .unwrap();
10 }
```

How is unwrap_or_else() in Option<T> defined?

```
1 impl<T> Option<T> {
2     pub fn unwrap_or_else<F>(self, f: F) → T
3     where
4         F: FnOnce() → T
5     {
6         match self {
7             Some(x) ⇒ x,
8             None ⇒ f(),
9         }
10    }
11 }
```

Fn Traits

- FnOnce: closures that can be called once
 - moves captured values out of its body
 - All closures implement at least this trait
- FnMut: closures that don't move captured values out of their body, but may mutate the captured values
- Fn: closures that don't move captured values out of their body, and don't mutate captured values, or capture nothing

Using the `sort_by_key` method with closures

```
1 #[derive(Debug)]
2 struct Rectangle {
3     width: u32,
4     height: u32,
5 }
6
7 fn main() {
8     let mut list = [
9         Rectangle { width: 10, height: 1 },
10        Rectangle { width: 3, height: 5 },
11        Rectangle { width: 7, height: 12 },
12    ];
13
14     // closure may be called multiple times
15     // so it takes an `FnMut` closure
16     list.sort_by_key(|r| r.width);
17     println!("{}list:#?{}");
18 }
```

Will this work?

```
1 let mut sort_operations = vec![];
2 let value = String::from("closure called");
3
4 list.sort_by_key(|r| {
5     sort_operations.push(value);
6     r.width
7});
```

No, since it moves `value` out of the closure by transferring ownership of `value` to `sort_operations`, and can only be called once — an `FnOnce` closure.

Will this work?

```
1 let mut num_sort_operations = 0;  
2  
3 list.sort_by_key(|r| {  
4     num_sort_operations += 1;  
5     r.width  
6 });
```

Yes, since it only captures a mutable reference to num_sort_operations — an FnMut closure.

Iterators

- Allows you to perform a task on a sequence of items
- Iterators are **lazy** — no effect until you call methods that consume the iterator to use it up

Using an iterator in a for loop

```
1 let v1 = vec![1, 2, 3];
2
3 let v1_iter = v1.iter();
4
5 for val in v1_iter {
6     println!("Got: {}", val);
7 }
```

The Iterator Trait and the next Method

```
1 pub trait Iterator {  
2     type Item;  
3  
4     fn next(&mut self) → Option<Self::Item>;  
5  
6     // methods with default implementations elided  
7 }
```

```
1 let v1 = vec![1, 2, 3];
2 let mut v1_iter = v1.iter();
3
4 assert_eq!(v1_iter.next(), Some(&1));
5 assert_eq!(v1_iter.next(), Some(&2));
6 assert_eq!(v1_iter.next(), Some(&3));
7 assert_eq!(v1_iter.next(), None);
```

`iter()` vs. `into_iter()` vs. `iter_mut()`

- With `iter()`, `next()` returns immutable references to values in the vector
- To create an iterator that takes ownership, call `into_iter()`
- To iterate over mutable references, call `iter_mut()`

Methods that Consume the Iterator

Methods, such as `sum()`, that call `next()` are called **consuming adaptors**, because calling them uses up the iterator

```
1 fn iterator_sum() {  
2     let v1 = vec![1, 2, 3];  
3     let v1_iter = v1.iter();  
4  
5     let total: i32 = v1_iter.sum();  
6     assert_eq!(total, 6);  
7 }
```

Methods that Produce Other Iterators

- **Iterator adaptors** don't consume the iterator
- Instead, they produce different iterators by changing some aspect of the original iterator

Calling the iterator adaptor `map()` to create a new iterator

```
1 let v1: Vec<i32> = vec![1, 2, 3];  
2  
3 v1.iter().map(|x| x + 1);
```

note: iterators are lazy and do nothing unless consumed

We need to consume the iterator — they are **lazy** — and the closure never gets called!

```
1 let v1: Vec<i32> = vec![1, 2, 3];
2
3 // collect() consumes the iterator
4 let v2: Vec<_> = v1.iter().map(|x| x + 1).collect();
5
6 assert_eq!(v2, vec![2, 3, 4]);
```

Using Closures that Capture Their Environment

```
1 #[derive(PartialEq, Debug)]
2 struct Shoe {
3     size: u32,
4     style: String,
5 }
6
7 fn shoes_in_size(shoes: Vec<Shoe>, shoe_size: u32) -> Vec<Shoe> {
8     // filter() takes a closure that captures `shoe_size`
9     shoes.into_iter().filter(|s| s.size == shoe_size).collect()
10 }
```

Required Additional Reading

The Rust Programming Language, Chapter 13.1, 13.2