

Lifetimes

Performant Software Systems with Rust — Lecture 9

Baochun Li, Professor

Department of Electrical and Computer Engineering
University of Toronto

Lifetimes — a First Cut

- Every **reference** has a lifetime
 - Typically (and in early versions of Rust), it's the scope for which the reference is valid
 - We will see variations soon
- Just like type inference, lifetimes are **inferred** by the compiler in most cases
- But also like type annotation, we must annotate lifetimes when inference is not possible

Consider This Example

```
1 fn main() {  
2     let r;  
3     {  
4         let x = 5;  
5         r = &x;  
6     }  
7  
8     println!("r: {}", r);  
9 }
```

What will happen at compile-time?

```
error[E0597]: `x` does not live long enough
→ src/main.rs:6:13
5      let x = 5;
6          - binding `x` declared here
7          r = &x;
8              ^` borrowed value does not live long enough
9      }
10     - `x` dropped here while still borrowed
11
12     println!("r: {}", r);
13             --- borrow later used here
```

The Borrow Checker

```
1 fn main() {  
2     let r; // -----+--- 'a  
3     {  
4         let x = 5; // +--- 'b /  
5         r = &x; // | /  
6     } // x goes out of scope // -+ /  
7     // /  
8     println!("r: {}", r); // /  
9 }
```

- This code is rejected at compile-time because `x`'s lifetime, '`b`, is not as long as `r`'s lifetime, '`a`
- Or, as the compiler says, *x does not live long enough!*

Lifetime Annotations in Functions

```
1 fn main() {
2     let s1 = String::from("abcd");
3     let s2 = "xyz";
4
5     // `longest()` takes string slices as we don't want it
6     // to take ownership
7     let result = longest(s1.as_str(), s2);
8     println!("The longest string is {}", result);
9 }
```

Implementing longest()

```
1 fn longest(x: &str, y: &str) → &str {  
2     if x.len() > y.len() {  
3         x  
4     } else {  
5         y  
6     }  
7 }
```

```
error[E0106]: missing lifetime specifier
→ src/main.rs:1:33
1 fn longest(x: &str, y: &str) → &str {
               -----           ^ expected named lifetime parameter
= help: this function's return type contains a borrowed value, but the
  signature does not say whether it is borrowed from `x` or `y`
help: consider introducing a named lifetime parameter
1 fn longest<'a>(x: &'a str, y: &'a str) → &'a str {
```

Lifetime Annotations

- We need to define the **relationship between the references** using **lifetime annotations**, so the borrow checker can perform its analysis
- **Lifetime annotations** don't change how long any of the references live — they are just hints to the **borrow checker**

```
1 &i32          // a reference
2 &mut i32      // a mutable reference
3 &'a i32       // a reference with an explicit lifetime
4 &'a mut i32   // a mutable reference with an explicit lifetime
```

Revisiting longest()

```
1 fn longest<'a>(x: &'a str, y: &'a str) → &'a str {  
2 // the returned reference will live as long as 'a  
3 // or, the returned reference will be valid as long as both the  
4 // parameters are valid  
5 // or, the returned reference cannot outlive either x or y  
6 // or, the lifetime of the returned reference is the same as the  
7 // smaller of the lifetimes of the two references passed in  
8 ...  
9 }
```

The Borrow Checker: Working with Annotated Lifetimes

```
1 fn main() {
2     let s1 = String::from("long string is long");
3
4     {
5         let s2 = String::from("xyz");
6         let result = longest(s1.as_str(), s2.as_str());
7         println!("The longest string is {result}");
8     }
9 }
```

The longest string is long string is long
()

The Borrow Checker: Working with Annotated Lifetimes

```
1 fn main() {
2     let s1 = String::from("long string is long");
3     let result;
4     {
5         let s2 = String::from("xyz");
6         result = longest(s1.as_str(), s2.as_str());
7     }
8
9     println!("The longest string is {}", result);
10 }
```

```
error[E0597]: `s2` does not live long enough
→ src/main.rs:14:39
```

```
13     let s2 = String::from("xyz");
      -- binding `s2` declared here
14     result = longest(s1.as_str(), s2.as_str());
                           ^^^ borrowed value does not live long
15 }
   - `s2` dropped here while still borrowed
16 println!("The longest string is {result}");
                           ----- borrow later used here
```

Lifetime Annotations in Structs

If a **struct** holds references, we need a lifetime annotation for each reference

```
1 struct ImportantExcerpt<'a> {
2     // an instance of `ImportantExcerpt` can't outlive the reference
3     // it holds in `part`
4     part: &'a str,
5 }
6
7 fn main() {
8     let novel = String::from("Rust vs. C++");
9     let first_sentence = novel.split('.').next().unwrap();
10    let i = ImportantExcerpt {
11        part: first_sentence,
12    };
13    println!("{}", i.part);
14 }
```

Lifetime Elision

- Each input parameter gets its own lifetime
- If there is exactly one input lifetime parameter, its lifetime is assigned to all output lifetime parameters
- If there are multiple input lifetime parameters, but one of them is `&self` or `&mut self` because this is a method, the lifetime of `self` is assigned to all output lifetime parameters

Lifetime Elision: Example

```
1 fn first_word(s: &str) → &str {...}
```



```
1 fn first_word<'a>(s: &'a str) → &'a str {...}
```

Lifetime Annotations in Method Definitions

```
1 struct ImportantExcerpt<'a> {
2     part: &'a str,
3 }
4
5 impl<'a> ImportantExcerpt<'a> {
6     fn announce_and_return_part(&self, announcement: &str) → &str {}
7 }
```



```
1 impl<'a, 'b> ImportantExcerpt<'a, 'b> {
2     fn announce_and_return_part(&'a self, announcement: &'b str)
3             → &'a str {}
4 }
```

The '`static`' Lifetime

- A special lifetime that tells the compiler that the reference can live for the entire duration of the program
- All string literals have the '`static`' lifetime, as they are in **static memory** (or the `data` segment)
- Watch out on following the compiler's suggestions — do not use '`static`' lifetimes if you don't know what you are doing

```
1 let s: &'static str = "I have a static lifetime.>";
```

Anonymous Lifetimes

- The compiler should try to infer the lifetime annotation by itself
- It is typically for simplifying the grammar when
 - writing `impl` blocks
 - returning structs and enums with annotated lifetimes

```
1 struct ImportantExcerpt<'a> {
2     // an instance of `ImportantExcerpt` can't outlive the reference
3     // it holds in `part`
4     part: &'a str,
5 }
6
7 // impl<'a> ImportantExcerpt<'a> {
8 impl ImportantExcerpt<'_> {
9     fn print(&self) {
10         println!("{}", self.part);
11     }
12 }
```

```
1 // impl<'a> ImportantExcerpt<'a> {
2 impl ImportantExcerpt<'_> {
3     fn print(&self) {
4         println!("{}", self.part);
5     }
6
7     // fn get_part<'a>(&self) -> &'a str {
8     fn get_part(&self) -> &'_ str {
9         self.part
10    }
11 }
```

Lifetimes **from scratch** again: an **in-depth** coverage

Let's revisit our first example:

```
1 fn main() {  
2     let r; // -----+-- 'a  
3     {  
4         let x = 5; // +-- 'b /  
5         r = &x; // | /  
6     } // x goes out of scope // +- /  
7     // /  
8     println!("r: {}", r); // /  
9 }
```

```
1 error[E0597]: `x` does not live long enough
```

There's just one problem, though.

Let's consider the following revised code:

```
1 fn main() {  
2     let r;           // -----+--- 'a  
3     {  
4         let x = 5;    // +--- 'b  |  
5         r = &x;       // |  |  
6     } // x goes out of scope // +- |  
7     // |  
8 }
```

Now the code compiles successfully. But why?

Lifetime vs. Scope: Revised Code

```
1 fn main() {  
2     let r;  
3     {  
4         let x = 5;           // x is not a reference, no lifetime!  
5         r = &x;             // -----+-- 'b  
6     } // x goes out of scope  
7 }
```

Lifetime vs. Scope: Original Example

```
1 fn main() {  
2     let r;  
3     {  
4         let x = 5;          // x is not a reference, no lifetime!  
5         r = &x;           // -----+-- 'b  
6     } // x goes out of scope // -+      /  
7     println!("r: {}", r);  //           /  
8 }
```

Let's Consider Another Example

```
1 fn main() {  
2     let foo = 1;  
3     let mut r;  
4     {  
5         let x = 5;  
6         r = &x;  
7  
8         println!("r: {}", r);  
9     }  
10    r = &foo;  
11    println!("r: {}", r);  
12 }  
13 }
```

Now the code compiles successfully. But why?

Let's Look At Lifetimes Again

```
1 fn main() {
2     let foo = 1;
3     let mut r;
4     {
5         let x = 5;
6         r = &x;           // -----+-- 'b
7
8         println!("r: {}", r); // /-----+
9     }                   // 'b is not alive
10    // 'b is not alive
11    r = &foo;          // -----+-- 'b
12    println!("r: {}", r); // -----+
13 }
```

Let's Make Another Revision

```
1 fn main() {
2     let foo = 1;
3     let r;
4     {
5         let x = 5;
6         r = &foo;           // -----+-- 'b
7         //-----+
8         println!("r: {}", r); //-----+
9     }
10    //-----+
11    println!("r: {}", r); //-----+
12 }
```

The revised code also compiles successfully. But why?

Let's Look At a Third Example

```
1 use rand::Rng;
2
3 fn main() {
4     let mut rng = rand::rng();
5     let mut x = String::from("Hello");
6     let random_float: f64 = rng.random();
7     let r = &x; // -----+-- 'b, immutable borrow on x
8     //      /
9     if random_float > 0.5 {
10         // 'b is not alive
11         x.push_str(" World!"); // 'b is not alive, mutable borrow on x
12     } else {
13         //      /
14         println!("{}"); // -----+-- 'b
15     }
16 }
```

This example compiles and runs successfully!

But What If We Add One Line of Code?

```
1 use rand::Rng;
2
3 fn main() {
4     let mut rng = rand::rng();
5     let mut x = String::from("Hello");
6     let random_float: f64 = rng.random();
7     let r = &x;                                // -----+ 'b, immutable borrow on x
8
9     if random_float > 0.5 {                  // 'b is alive
10        x.push_str(" World!");             // 'b is alive, mutable borrow on x
11        println!("{}");                   // immutable borrow!
12    } else {                                // -----+
13        println!("{}");                   // -----+ 'b
14    }
15 }
```

```
1 error[E0502]: cannot borrow `x` as mutable because it is also
2 borrowed as immutable
3     → src/main.rs:11:9
4
5 7         let r = &x;
6             -- immutable borrow occurs here
7 ...
8 11         x.push_str(" World!");
9             ^^^^^^^^^^^^^^^^^^ mutable borrow occurs here
10 12         println!("{}");           --- immutable borrow later used here
11
```

The data flow model

— Chapter 1, Foundations, in Rust for Rustaceans, Jon Gjengset

References

- There are two kinds of references
 - **Shared reference:** &
 - **Mutable reference:** &mut
- All references obey the following rules
 - A reference cannot outlive its referent
 - A mutable reference cannot be **aliased**
- But what do we mean by **aliasing**?

Aliasing

Variables and pointers **alias** if they refer to overlapping regions of memory

Consider the following code:

```
1 fn compute(input: &u32, output: &mut u32) {  
2     if *input > 10 {  
3         *output = 1;  
4     }  
5     if *input > 5 {  
6         *output *= 2;  
7     }  
8     // `*output` will be `2` if `input > 10`  
9 }
```

Can the Rust compiler optimize the previous code to the following?

```
1 fn compute(input: &u32, output: &mut u32) {  
2     let cached_input = *input; // keep `*input` in a register  
3     if cached_input > 10 {  
4         *output = 2;  
5     } else if cached_input > 5 {  
6         *output *= 2;  
7     }  
8 }
```

Alias Analysis Helps

- We used the fact that `&mut u32` can't be aliased to prove that writes to `*output` can't possibly affect `*input`
- This lets us cache `*input` in a register, eliminating a read
- **Alias analysis** lets the compiler perform useful optimizations!

But should we be concerned with aliasing in the following modified code?

```
1 fn compute(input: &u32, output: &mut u32) {  
2     let mut temp = *output;  
3     if *input > 10 {  
4         temp = 1;  
5     }  
6     if *input > 5 {  
7         temp *= 2;  
8     }  
9     *output = temp;  
10 }
```

No. `input` doesn't alias `temp`, because the value of a local variable can't be aliased by things that existed before it was declared.

The definition of **alias** in Rust needs to involve some notion of **liveness** and **mutation** — we don't actually care if aliasing occurs if there aren't any actual writes to memory happening

Now Let's Revisit Lifetimes

- A **Lifetime** involves **named regions of code** that a reference must be valid — **alive** — for
- A lifetime corresponds to a **path of execution**, with potential holes in them
- In most cases, a reference's lifetime coincides to its scope
- Most lifetimes in Rust are inferred by the compiler

```
1 let mut data = vec![1, 2, 3];
2 let x = &data[0];
3 data.push(4);
4 println!("{}", x);
```

And the compiler sees

```
1 'a: {
2     let mut data: Vec<i32> = vec![1, 2, 3];
3     'b: {
4         // 'b is as big as we need this borrow to be
5         // (just need to get to `println!`)
6         let x: &'b i32 = Index::index▷◁'b>(&'b data, 0);
7         'c: {
8             // Temporary scope because we don't need the
9             // &mut to last any longer
10            Vec::push(&'c mut data, 4);
11        }
12        println!("{}", x);
13    }
14 }
```

```
error[E0502]: cannot borrow `data` as mutable because it is also
borrowed as immutable
```

Recommended two-hour video: Crust of Rust: Lifetime Annotations, Jon Gjengset

Required Additional Reading

The Rust Programming Language, Chapter 10.3

The Rustonomicon, Chapter 3.1 - 3.3