

Ownership and the String Type

Performant Software Systems with Rust — Lecture 4

Baochun Li, Professor

Department of Electrical and Computer Engineering
University of Toronto

The Dreaded error[E0382]

```
1 error[E0382]: borrow of moved value
```

Let's start from the very beginning

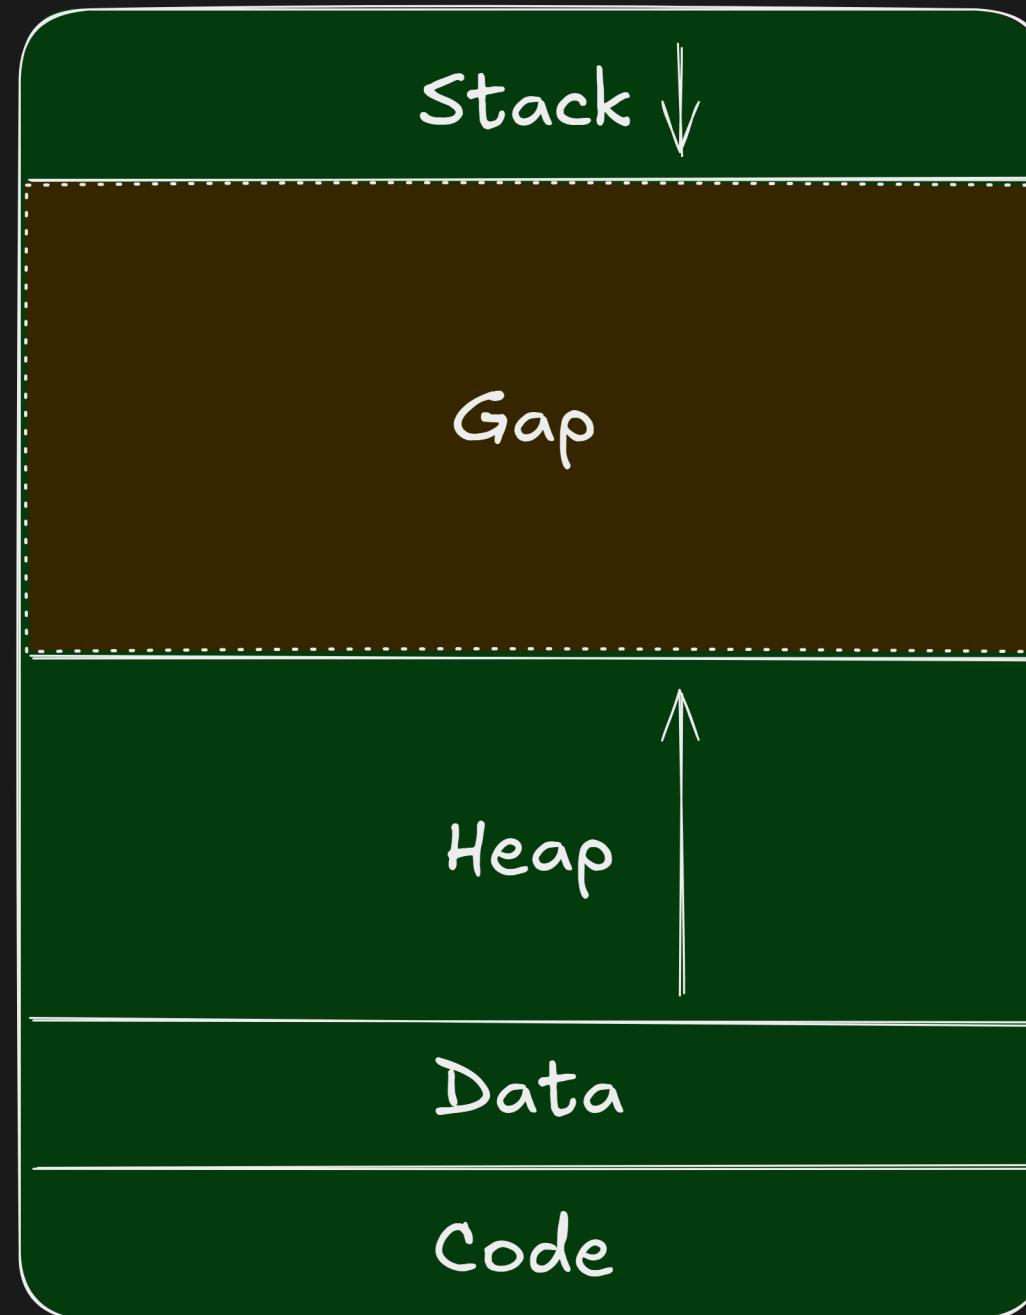
The Stack

- Variables and parameters on the stack are **local** to the function scope they are declared in
- Function calls proceed in a **last-in, first-out** order
 - When we call a function, we **push** its stack frame onto the stack
 - When we return from it, we **pop** its stack frame off
- The stack is **fast** — we only need to move the **stack pointer**

The Heap

- The heap is where **dynamic memory** is allocated — we need to explicitly request and release memory
 - A request (with `malloc()` in C or `new` in C++) returns a pointer to allocated memory on the heap
 - This pointer can be stored in a variable on the stack
 - A release (with `free()` in C or `delete` in C++) deallocates the memory at a pointer
- Slower than the stack because we need to **search** for a block of memory that is large enough

The stack grows **downwards** to **lower** memory addresses, and the heap grows **upwards** to **higher** memory addresses



Ownership Rules

- Each value in Rust has a variable that is its **owner**
- There can only be **one** owner at a time
- When the owner goes out of scope, the value will be **dropped**

Revisiting Variable Scope

```
1 { // s is not valid here, it's not yet declared
2     let s = "hello";
3     // s comes into scope, it is valid from this point forward
4     println!("{}");
5     // s remains valid until it goes out of scope
6 } // this scope is now over, s is no longer valid
```

hello

Fun fact: This runtime output is produced by running the code directly when compiling this presentation, with a Jupyter Kernel for Rust.

To continue to study **ownership**, we need a data type that is stored on the heap

The String Type

- We've seen string literals before, but they are **immutable**
- What if we want to store a string that is **unknown** at compile time?

```
1 // a mutable String allocated on the heap
2 let mut s = String::from("hello");
3 s.push_str(", world!"); // appends a literal to an existing String
4 println!("{}");
```

hello, world!

Allocating and Releasing Strings on the Heap

- We allocated memory for the String `s` manually using `String::from`
- When the variable `s` goes out of scope, the memory is automatically released

Rust vs. C/C++

- In Rust, there is no need to manually call `free` or `delete` explicitly
 - Wastes memory if we do it too late (or forget to do it — **memory leaks**)
 - Invalid memory access if we do it too early
 - Double free if we do it twice, corrupting the memory allocator

Rust vs. Go/JavaScript

- In Rust, there is no need to use a **garbage collector** to clean up unused memory
- Allocated memory is **automatically** released once the variable that owns it goes out of scope

```
1 { // s is not valid here, it's not yet declared
2     let mut s = String::from("hello");
3     // s comes into scope, it is valid from this point forward
4     s.push_str(", world!");
5     // s remains valid until it goes out of scope
6     println!("{}");
7 } // this scope is now over, s is no longer valid and is `dropped'
```

Resource Acquisition Is Initialization (RAII) in C++

- Rust's idea is not new — it is inspired by RAII in C++
- With RAII, resources are acquired in the constructor and released in the destructor of an object
- This is why C++ has **smart pointers** like `std::unique_ptr` and `std::shared_ptr`
- The main difference is
 - Rust doesn't have `new` and `delete`
 - Rust enforces the ownership rules at compile time

Difference Between Scalar Types and Strings

```
1 let x = 5; // bind the value 5 to x
2 let y = x; // make a copy of the value in x and bind it to y
3 println!("{} , {}");
```

5, 5

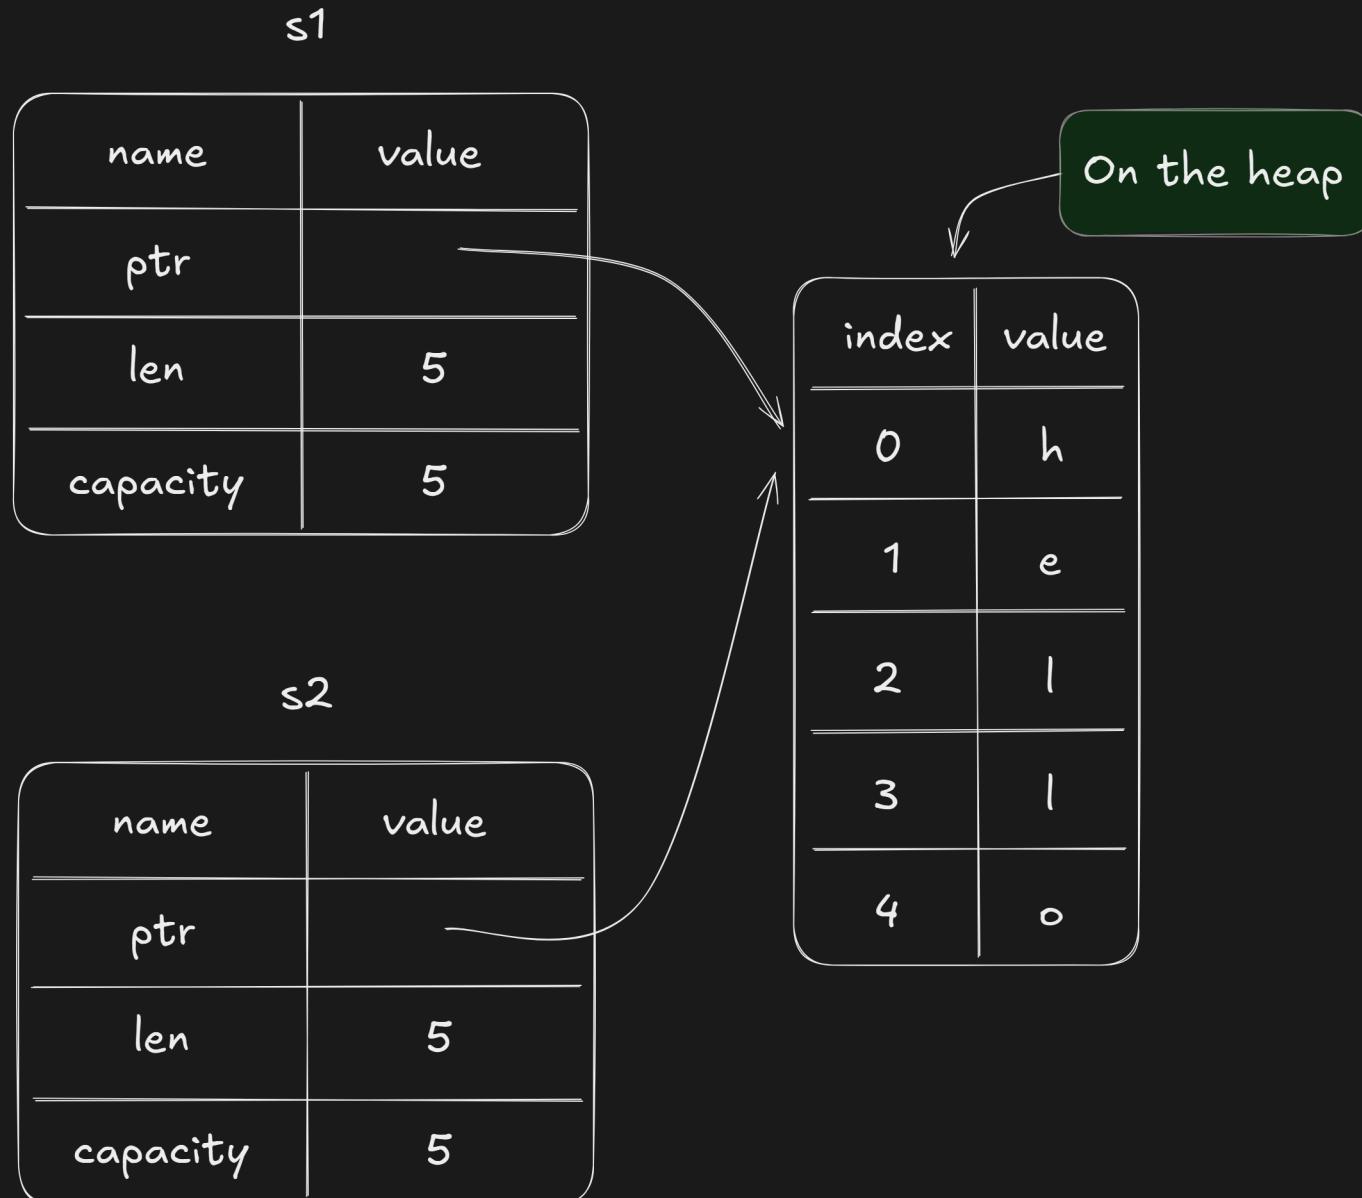
```
1 let s1 = String::from("hello");
2 let s2 = s1;
3 println!("{} , {}");
```

?

```
error[E0382]: borrow of moved value: `s1`
--> main.rs:4:15
|
2 |     let s1 = String::from("hello");
|         -- move occurs because `s1` has type `String`, which does not implement the `Copy` trait
3 |     let s2 = s1;
|             -- value moved here
4 |     println!("{}{}, {}", s1, s2);
|                     ^^^^ value borrowed here after move
|
= note: this error originates in the macro `$crate::format_args_nl` which comes from the expansion of the macro `println` (in Nightly builds, run with -Z macro-backtrace for more info)
help: consider cloning the value if the performance cost is acceptable
|
3 |     let s2 = s1.clone();
|             ++++++++
|
error: aborting due to 1 previous error

For more information about this error, try `rustc --explain E0382`.
```

What Happens When You Copy a String?



Back to Our Example

```
1 {
2     let s1 = String::from("hello");
3     let s2 = s1;
4     println!("{}{}, {}", s1, s2);
5 } // s1 and s2 go out of scope at the same time
6 // but should both of them be dropped? --- ``double free'' error!
```

Making a Move

We should make a **shallow copy** and then **invalidate** the original variable copied — which is called a **move**

```
1 {
2     let s1 = String::from("hello");
3     let s2 = s1; // s1 is moved into s2, and is no longer valid
4     println!("{}{}, {}", s1, s2); // error: borrow of moved value: `s1`
5 } // s1 and s2 go out of scope at the same time
6 // but s1 does not need to be freed, as it is no longer valid
```

What happens with this code?

```
1 let mut s = String::from("hello");
2 s = String::from("ahoy");
3
4 println!("{} , world!", s);
```

Making a Clone

```
1 {
2     let s1 = String::from("hello");
3     let s2 = s1.clone(); // s2 has a deep copy of s1
4     println!("{} , {}"); // no more errors
5 } // s1 and s2 go out of scope at the same time, are are both freed
```

Back to Our Example Again

```
1  let x = 5; // bind the value 5 to x
2  let y = x; // make a copy of the value in x and bind it to y
3  println!("{}{}, {}", x, y);
```

5, 5

We haven't called `clone()` — why do we make a copy of `x` here?

- `x` is an integer on the stack only, and such stack-only data is not expensive to copy
- There is no reason to invalidate `x` and move it into `y` here
- `Copy` trait is implemented for all types that are stored on the stack only, just like integers
 - If a type implements the `Copy` trait, variables that use it do not move, but rather are trivially copied
 - If a type implements the `Drop` trait, it cannot implement the `Copy` trait

Passing a Value to a Function Can Transfer Ownership

```
1 fn main() {
2     let s = String::from("hello"); // s comes into scope
3     takes_ownership(s); // s's value moves into the function,
4     // s is no longer valid here
5     let x = 5; // x comes into scope
6     makes_copy(x); // x is an integer that implements the Copy trait,
7                     // copies into the function and okay to use later
8 } // both s and x goes out of scope
9 // Because s's value was moved, nothing special happens
```

```
1 fn takes_ownership(s: String) {  
2     println!("{}");  
3 } // s goes out of scope and is dropped  
4  
5 fn makes_copy(n: i32) {  
6     println!("{}");  
7 } // n goes out of scope, and nothing special happens
```

Returning a Value from a Function Can Also Transfer Ownership

```
1 fn main() {  
2     let s1 = gives_ownership(); // moves its return value into s1  
3     let s2 = String::from("hello");  
4     let s3 = takes_and_gives_back(s2);  
5     // s2 moved into the function; its return value moved into s3  
6 } // s1 and s3 are dropped, s2 was moved and nothing happens  
7  
8 fn gives_ownership() → String {  
9     let s = String::from("yours");  
10    s // moves `s` out of the function  
11 }  
12  
13 fn takes_and_gives_back(s: String) → String {  
14     s // moves into the function and then out of the function  
15 }
```

Allowing a Function to Use a Value Without Taking Ownership

```
1 fn main() {  
2     let s1 = String::from("hello");  
3     let (s1, len) = calculate_length(s1);  
4     println!("The length of '{s1}' is {len}.");  
5 }  
6  
7 fn calculate_length(s: String) → (String, usize) {  
8     (s, s.len())  
9 }
```

Any issues you can see here?

Demo: error[E0382]: borrow of moved value

Allowing a Function to Use a Value Without Taking Ownership

```
1 fn main() {
2     let s1 = String::from("hello");
3     let (s1, len) = calculate_length(s1);
4     println!("The length of '{s1}' is {len}.");
5 }
6
7 fn calculate_length(s: String) → (String, usize) {
8     let length = s.len();
9     (s, length)
10 }
```

The length of 'hello' is 5.

References and Borrowing

- **References** allow you to **borrow** a value without taking ownership of it
 - Represented by &
 - Can **derefence** them with *
 - Unlike C pointers, they can never be null or “dangling”

References and Borrowing

```
1 fn main() {  
2     let s1 = String::from("hello");  
3     let len = calculate_length(&s1); // pass a reference to s1  
4     println!("The length of '{s1}' is {len}.");  
5 }  
6  
7 fn calculate_length(s: &String) → usize {  
8     s.len()  
9 } // Here, s goes out of scope. But because s does not have ownership of wh  
10 // it refers to, the String is not dropped.
```

Trying to Modify What We Borrow

```
1 fn main() {
2     let s = String::from("hello");
3
4     change(&s);
5     println!("{}"), 
6 }
7
8 fn change(some_string: &String) {
9     some_string.push_str(", world");
10 }
```

Demo: error[E0596]: cannot borrow *some_string as mutable, as it is behind a & reference

Mutable References

- **Mutable references** allow you to change the value of a reference
 - Represented by `&mut`

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     change(&mut s); // pass a mutable reference to s  
5     println!("{}");  
6 }  
7  
8 fn change(some_string: &mut String) { // take a mutable reference  
9     some_string.push_str(", world");  
10 }
```

hello, world

Can we create **two** mutable references to
the same value?

```
1 let mut s = String::from("hello");
2
3 let r1 = &mut s;
4 let r2 = &mut s;
5
6 println!("{} , {}", r1, r2);
```

Demo: error[E0499]: cannot borrow `s` as mutable more than once at a time

Can we have **one** mutable reference and an immutable reference?

```
1 let mut s = String::from("hello");
2
3 let r1 = &s; // no problem
4 let r2 = &s; // no problem
5 let r3 = &mut s; // error[E0502]
6
7 println!("{} , {} , and {}", r1, r2, r3);
```

Demo: error[E0502]: cannot borrow `s` as mutable because it is also borrowed as immutable

Mutable References: House Rules

- Cannot borrow as a **mutable reference** more than once
- Cannot have a **mutable reference** while also having an **immutable reference**
- Can have **multiple immutable references** at the same time
- Summary: Can either have **one mutable reference** or **multiple immutable references**, but not both
- But why?

Data Races: The Most Elusive Bugs in Concurrent Software

- Two or more pointers access the same data at the same time
- At least one of the pointers is being used to write to the data
- There's no mechanism being used to synchronize access to the data

Data races in Rust are **compile-time** errors

Fixing the Compile-Time Error

```
1 let mut s = String::from("hello");
2
3 let r1 = &s; // no problem
4 let r2 = &s; // no problem
5 println!("{} and {}", r1, r2);
6 // variables r1 and r2 will not be used after this point
7
8 let r3 = &mut s; // no problem
9 println!("{}");
```

hello and hello

hello

()

Dangling References

```
1 fn main() {  
2     let reference_to_nothing = dangle();  
3 }  
4  
5 fn dangle() → &String { // dangle returns a reference to a String  
6     let s = String::from("hello"); // s is a new String  
7     &s // we return a reference to the String, s  
8 }  
9 // s goes out of scope, and is dropped.
```

Demo: error[E0106]: missing lifetime specifier

Fixing the Compile-Time Error

```
1 fn no_dangle() → String {  
2     let s = String::from("hello");  
3  
4     s // ownership is moved out, nothing is dropped  
5 }
```

References: House Rules

- References must always be valid
- At any given time, you can have either one mutable reference or any number of immutable references
- Data will not go out of scope before the reference to the data does

Problem: Write a function that takes a string of words, and returns the first word it finds in the string

```
1 fn first_word(s: &String) → ?
```

Returning an **index** of the end of the first word as an integer?

Let's try this idea.

```
1 fn first_word(s: &String) → usize {
2     let bytes = s.as_bytes(); // convert to an array of bytes
3
4     // .iter(): an iterator that returns each element
5     // .enumerate(): returns each element as a tuple
6     // (index, reference to element)
7     for (i, &item) in bytes.iter().enumerate() {
8         if item == b' ' {
9             return i;
10        }
11    }
12
13    s.len()
14 }
```

But there is no **guarantee that the index
returned will be **valid** in the future**

```
1 fn main() {
2     let mut s = String::from("hello world");
3
4     let word = first_word(&s); // word will get the value 5
5
6     s.clear(); // this empties s, making it equal to ""
7
8     // `word` still has the value 5 here, but there's no more
9     // string that we could meaningfully use the value 5 with.
10    // `word` is now invalid.
11 }
```

String Slices

- `&str`: **Immutable** references to a part of a `String`
- Created using a range within square brackets
 - `[starting_index..ending_index]`

```
1 let s = String::from("hello");
2 let slice = &s[0..2];
3 let slice = &s[..2];
4 let slice = &s[3..];
5 let slice = &s[..];
```

- Stored internally with the **starting reference** and **length** of the slice

String Literals as Slices

```
1 let s = "Hello, world!"; // the type of s is &str
```

Fixing Our Solution to `first_word()`

```
1 fn first_word(s: &String) → &str {
2     let bytes = s.as_bytes();
3
4     for (i, &item) in bytes.iter().enumerate() {
5         if item == b' ' {
6             return &s[0..i];
7         }
8     }
9
10    &s[..]
11 }
12
13 fn main() {
14     let mut s = String::from("hello world");
15     let word = first_word(&s);
16     s.clear(); // compile-time error[E0502]
17     println!("the first word is: {word}");
18 }
```

Demo: error[E0502]: cannot borrow `s` as mutable because it is also borrowed as immutable

String Slices as Parameters

Instead of `&String`, we can use `&str` as the parameter type

```
1 fn first_word(s: &str) → &str {
```

String Slices as Parameters

```
1 fn main() {
2     let my_string = String::from("hello world");
3
4     // works on slices of `String`s (partial or whole)
5     let word = first_word(&my_string[0..6]);
6     let word = first_word(&my_string[..]);
7     // also works on references to `String`s, which are
8     // equivalent to whole slices of `String`s
9     let word = first_word(&my_string);
10
11    let my_string_literal = "hello world";
12    // works on slices of string literals (partial or whole)
13    let word = first_word(&my_string_literal[0..6]);
14    let word = first_word(&my_string_literal[..]);
15
16    // Because string literals are string slices already,
17    // this works too, without the slice syntax!
18    let word = first_word(my_string_literal);
```

Creating Slices of Collection Types

```
1 let a = [1, 2, 3, 4, 5];
2 let slice = &a[1..3]; // `slice` has a type `&[i32]`
3 println!("{}", slice == &[2, 3]);
```

true

()

An **in-depth** look at the **String** type

What is a String?

- A **vector** of bytes

```
1 pub struct String {  
2     vec: Vec<u8>,  
3 }
```

What is a String?

- UTF-8-encoded and **growable**

```
1 let hello = String::from("السلام عليكم");  
2 let hello = String::from("Dobrý den");  
3 let hello = String::from("Hello");  
4 let hello = String::from("שלום");  
5 let hello = String::from("নমস্তে");  
6 let hello = String::from("こんにちは");  
7 let hello = String::from("안녕하세요");  
8 let hello = String::from("你好");  
9 let hello = String::from("Olá");  
10 let hello = String::from("Здравствуйте");  
11 let hello = String::from("Hola");
```

Creating and Initializing a New String

```
1 let mut s = String::new(); // creates a new, empty string
```

```
1 let s = String::from("hello");
```

```
1 let s = "hello".to_string();
```

Appending to a String

```
1 let mut s1 = String::from("foo");  
2 let s2 = "bar";  
3 s1.push_str(s2); // appends a string slice  
4 println!("s1 is {s1}, s2 is {s2}");
```

s1 is foobar, s2 is bar
()

```
1 let mut s = String::from("lo");  
2 s.push('l'); // appending a single character
```

Concatenating Strings with +

```
1 let s1 = String::from("Hello, ");
2 let s2 = String::from("world!");
3 // using deref coercion to turn `&s2` into `&s2[..]`
4 let s3 = s1 + &s2;
5 // s1 has been moved and is no longer valid
6 println!("{} , {}" , s3, s2);
```

```
Hello, world!, world!
()
```

```
1 fn add(self, s: &str) → String {
```

Concatenating Strings with the format! Macro

```
1 let s1 = String::from("tic");
2 let s2 = String::from("tac");
3 let s3 = String::from("toe");
4
5 let s = format!("{}-{}-{}", s1, s2, s3);
6 println!("{}", s);
```

```
tic-tac-toe
()
```

Indexing into Strings is Not Allowed

```
1 let s1 = String::from("hello");
2 let h = s1[0];
1 error[E0277]: the type `str` cannot be indexed by `'{integer}`
2 --> main.rs:3:16
3 |
4 3 |     let h = s1[0];
5 |             ^ string indices are ranges of `usize`
6 |
7 = help: the trait `SliceIndex<str>` is not implemented for `'{integer}`, w
8 = note: you can use `chars().nth()` or `bytes().nth()`
9      for more information, see chapter 8 in The Book: <https://doc.rust-lang.org>
10 = help: the trait `SliceIndex<[_]>` is implemented for `usize`
11 = help: for that trait implementation, expected `[_]`, found `str`
12 = note: required for `String` to implement `Index<{integer}>`
```

**Strings are encoded internally with UTF-8 —
a character can take one or two bytes**

Use Slices Instead

```
1 let hello = String::from("😊 😊 📜 ");
2 let s = &hello[0..4];
3 println!("{}", s);
```



()

Internal UTF-8 Encoding Values

```
1 let hello = String::from("😊 😊 📜 ");
2 let s = &hello[0..4];
3
4 for c in hello.chars() {
5     print!("{} ", c);
6 }
7
8 for b in s.bytes() {
9     println!("{} ", b);
10 }
```

😊 😊 📜 240

159

152

128

()

Required Additional Reading

The Rust Programming Language, Chapter 4, 8.2