

Smart Pointers

Performant Software Systems with Rust — Lecture 11

Baochun Li, Professor

Department of Electrical and Computer Engineering
University of Toronto

What is a Pointer?

- A **pointer** is a variable that contains an address in memory — it **points to** some other data
- In Rust, a pointer is a **reference**, indicated by &, which borrows the value it points to

What is a Smart Pointer?

- A **smart pointer** is a data structure that acts like a pointer, but also has additional metadata and capabilities
 - Implemented using `struct`, and implements `Deref` and `Drop` traits
 - In many cases, a smart pointer **owns** the data it points to
- Examples of smart pointers we used
 - `String`
 - `Vec<T>`

Box<T>

- Stores data on the heap, and the pointer on the stack
- **Zero** performance overhead, but no additional capabilities either

- Use Box<T> in three situations:
 - Transferring ownership with a large amount of data — avoids copying
 - **Trait objects** — will get to this later
 - Having a type without a known size at compile time, but want to use a value of this type in a context that requires an exact size

Using Box<T> to Store Data on the Heap

```
1 fn main() {  
2     // stores an i32 value on the heap using Box<T>  
3     let b = Box::new(5);  
4     println!("b = {}", b);  
5 }
```

b = 5

When a Box Goes Out Of Scope

- It will be deallocated
- The deallocation happens both for the **box** (stored on the stack) and the data it points to (stored on the **heap**)

Trait Objects

Recall that we can use an `enum` to store different types of data in each cell, while still having a vector of these cells

```
1 enum SpreadsheetCell {  
2     Int(i32),  
3     Float(f64),  
4     Text(String),  
5 }  
6  
7 let row = vec![  
8     SpreadsheetCell::Int(3),  
9     SpreadsheetCell::Text(String::from("blue")),  
10    SpreadsheetCell::Float(10.12),  
11];
```

There Is a Problem with Our Approach

- The types of cells in the vector are **fixed**
- But sometimes we are writing a library (crate) to be used by others, and want our library user to be able to extend the set of types
- For example, in a graphical UI, we wish the vector to store all the objects that can draw themselves
 - sounds like **abstract classes** in object-oriented programming!

Towards Using Trait Objects

- Let's say we wish to implement a **vector** that contains animals that can make a sound in our library
- But the types of these animals are to be defined by the users of our library
- We first define a **trait**

```
1 trait Animal {  
2     fn make_sound(&self);  
3 }
```

Towards Using Trait Objects

- A user of our library defines two concrete types: Dog and Cat
- These types are of different sizes

```
1 struct Dog {  
2     name: String  
3 }  
4  
5 struct Cat {  
6     lives: u8  
7 }  
8  
9 impl Animal for Dog {  
10    fn make_sound(&self) {  
11        println!("Woof!");  
12    }  
13 }  
14  
15 impl Animal for Cat {  
16    fn make_sound(&self) {  
17        println!("Meow!");  
18    }  
19 }
```

Placing Trait Objects in Vectors

```
1 fn main() {  
2     // trait objects!  
3     let animals: Vec<dyn Animal> = vec![  
4         Dog { name: String::from("Rover") },  
5         Cat { lives: 9 }  
6     ];  
7  
8     for animal in animals {  
9         animal.make_sound();  
10    }  
11 }
```

Will it compile successfully?

Live Demo

Correct Solution

```
1 fn main() {
2     let animals: Vec<Box<dyn Animal>> = vec![
3         Box::new(Dog { name: String::from("Rover") }),
4         Box::new(Cat { lives: 9 })
5     ];
6
7     // Use iter() to borrow each element
8     for animal in animals.iter() {
9         animal.make_sound();
10    }
11 }
```

Why Boxing Trait Objects?

- Trait objects are **dynamically sized types** (DSTs) — we don't know their sizes at compile-time
- But the `T` in `Vec<T>` must implement the `Sized` trait
 - all types with known sizes automatically implement the `Sized` trait
- Boxed trait objects have **known sizes** — we know the size of a pointer!

Generic Types Implement Sized

```
1 fn generic<T>(t: T) {}
```

is the same as

```
1 fn generic<T: Sized>(t: T) {}
```

unless we explicitly opt out:

```
1 fn generic<T: ?Sized>(t: &T) {}
```

Dynamically Sized Types and Wide Pointers

- `str` is a **dynamically sized type**
 - since we don't know the size of a string slice
- `&str`, however, has a **known size**
 - it is a **wide pointer** (also called a **fat pointer**)
 - contains the address of `str` and its **length**

Wide Pointers for Slices

- Slices, such as `str` or `[T]`, is simply a view into some continuous data, such as a vector
- A **wide pointer** to a slice contains the address and the number of elements

Wide Pointers for Trait Objects

- A **wide pointer** to a trait object, such as `dyn Animal`, consists of a data pointer and a `vtable` pointer
 - the data pointer addresses the data (of some unknown type `T`) that the trait object is storing
 - the `vtable` is a `struct` of function pointers, pointing to the concrete piece of machine code for each method

Back to Box<T>

- Implements the `Deref` trait, allowing its values to be treated like **references**
- Implements the `Drop` trait, allowing the memory it boxes to be deallocated

Box<T> Implements the Deref Trait

```
1 fn main() {  
2     let x = 5;  
3     let y = Box::new(x);  
4  
5     assert_eq!(5, x);  
6     assert_eq!(5, *y);  
7 }
```

Defining Our Own Simple Box

```
1 // MyBox is simply a tuple struct with one element of type T
2 #[derive(Debug)]
3 struct MyBox<T>(T);
4
5 impl<T> MyBox<T> {
6     fn new(x: T) -> MyBox<T> {
7         MyBox(x)
8     }
9 }
```

Can it be dereferenced?

```
1 fn main() {
2     let y = MyBox::new(5);
3     assert_eq!(5, *y);
4 }
```

Implementing the Deref Trait

```
1 use std::ops::Deref;  
2  
3 impl<T> Deref for MyBox<T> {  
4     type Target = T; // uses an associated type  
5  
6     fn deref(&self) → &Self::Target {  
7         &self.0  
8     } // *y → *(y.deref())  
9 }
```

Deref Coercion

```
1 fn main() {  
2     let m = MyBox::new(String::from("Rust"));  
3     let hello = |name: &str| → println!("Hello, {name}!");  
4     hello(&m); // equivalent to hello(&(*m)[..])  
5 }
```

Hello, Rust!

Deref Coercion and the DerefMut Trait

- From `&T` to `&U` when `T: Deref<Target=U>`
- From `&mut T` to `&mut U` when `T: DerefMut<Target=U>`
- From `&mut T` to `&U` when `T: Deref<Target=U>`

The Drop Trait

By implementing the `drop` method (that takes `&mut self`) in the `Drop` trait, you specify the code to run when a value goes out of scope

```
1 impl<T> Drop for MyBox<T> {
2     fn drop(&mut self) {
3         println!("Dropping MyBox<T>!");
4     }
5 }
```

Dropping a Value Early by with std::mem::drop

```
1 fn main() {
2     let y = MyBox::new(5);
3     assert_eq!(5, *y);
4
5     let m = MyBox::new(String::from("Rust"));
6     let hello = |name: &str| println!("Hello, {}!", name);
7     drop(y);
8     hello(&m);
9 }
```

Live Demo

Rc<T> — The Reference Counted Smart Pointer

- Allowing multiple ownership of the same data
- Read-only, no mutability
- Uses **reference counting** to ensure memory safety
- Single-threaded, otherwise use Arc<T>

```
1 use std::rc::Rc; // remember to import Rc
2
3 fn main() {
4     // Create a new reference-counted string
5     let first = Rc::new(String::from("Hello"));
6     println!("Rc after creation: {}", Rc::strong_count(&first));
7
8     {
9         // Create a second reference to the same data
10        let second = Rc::clone(&first);
11        println!("Rc after clone: {}", Rc::strong_count(&first));
12
13        // Both references can read the data
14        println!("First: {}", *first);
15        println!("Second: {}", *second);
16    } // second is dropped here
17
18 // Reference count decreases when second goes out of scope
```

RefCell<T> and the Interior Mutability Pattern

- **Interior mutability** is a design pattern in Rust that allows you to mutate data even when there are immutable references to that data
- We do this with RefCell<T>

RefCell<T>

- Single ownership, unlike Rc<T>
- Moves borrowing checks from compile-time to runtime
- Only use when compile-time borrowing rules are too restrictive
- **Panics** if borrowing rules are violated at runtime
- Often used with Rc for shared mutable state in single-threaded contexts
- Use Arc<Mutex<T>> instead of Rc<RefCell<T>> in multi-threaded contexts

```
1 use std::cell::RefCell;
2 use std::rc::Rc;
3
4 fn main() {
5     // Create a RefCell wrapped in an Rc for shared mutable access
6     let counter = Rc::new(RefCell::new(0));
7
8     // Create multiple references to the same data
9     let counter_ref1 = Rc::clone(&counter);
10    let counter_ref2 = Rc::clone(&counter);
11
12    // Modify the value through the first reference
13    *counter_ref1.borrow_mut() += 1;
14    // `borrow()` returns a Ref<T> which derefs to &T
15    println!("After first modification: {}", counter_ref1.borrow());
16
17    // Modify the value through the second reference
18    *counter_ref2.borrow_mut() += 2;
```

Live Demo

Required Additional Reading

The Rust Programming Language, Chapter 15.1 – 15.5, 18.2,
20.3

Rust for Rustaceans, Jon Gjenset, **Chapter 2 (Types)**